


Chapter 3: Processes





Chapter 3: Processes

- **Process Concept**
- **Operations on Processes**
- **Inter-process Communication(IPC)**





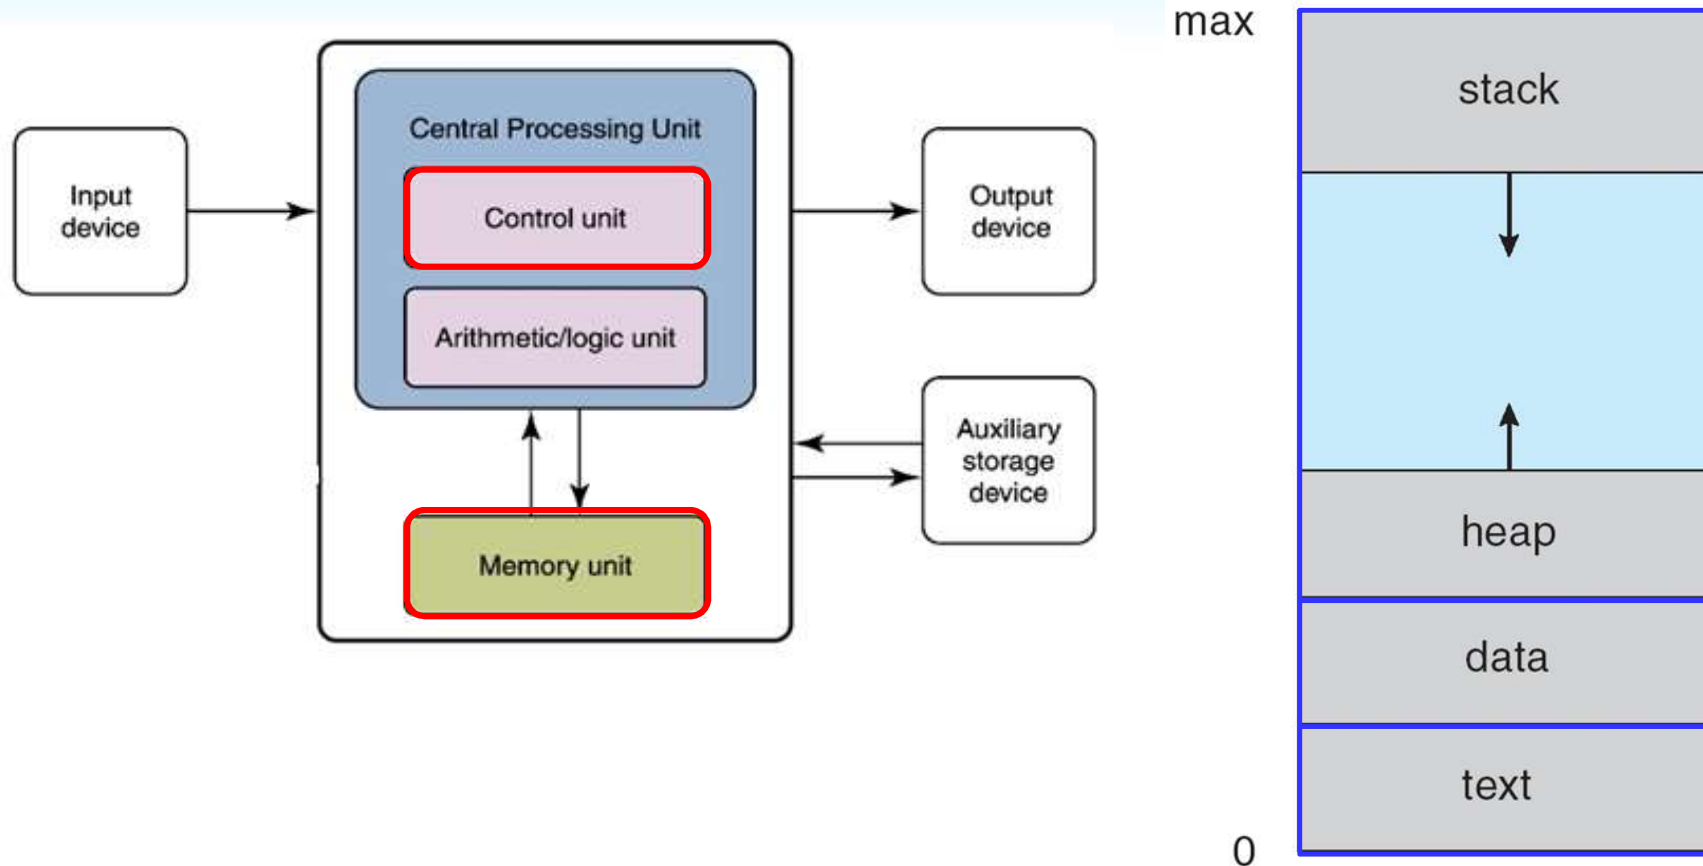
Process Concept

- **Process** – a **program in execution**; process execution must **progress in sequential fashion**
- Multiple parts
 - The program code, also called **text section**
 - **Data section** containing global variables, static local variables
 - **Stack** containing temporary data
 - ▶ Function parameters, return addresses, automatic local variables
 - **Heap** containing memory dynamically allocated during run time
 - Current activity including **program counter**, processor registers
- Program is **passive** entity stored on disk (executable file), **process is active**
 - **Program becomes process when executable file loaded into memory**
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be **several processes**
 - Consider multiple users executing the same program





Process in Memory





Stored Program Concept

High-level Language

```
int i, j, k;  
k = i + j;
```

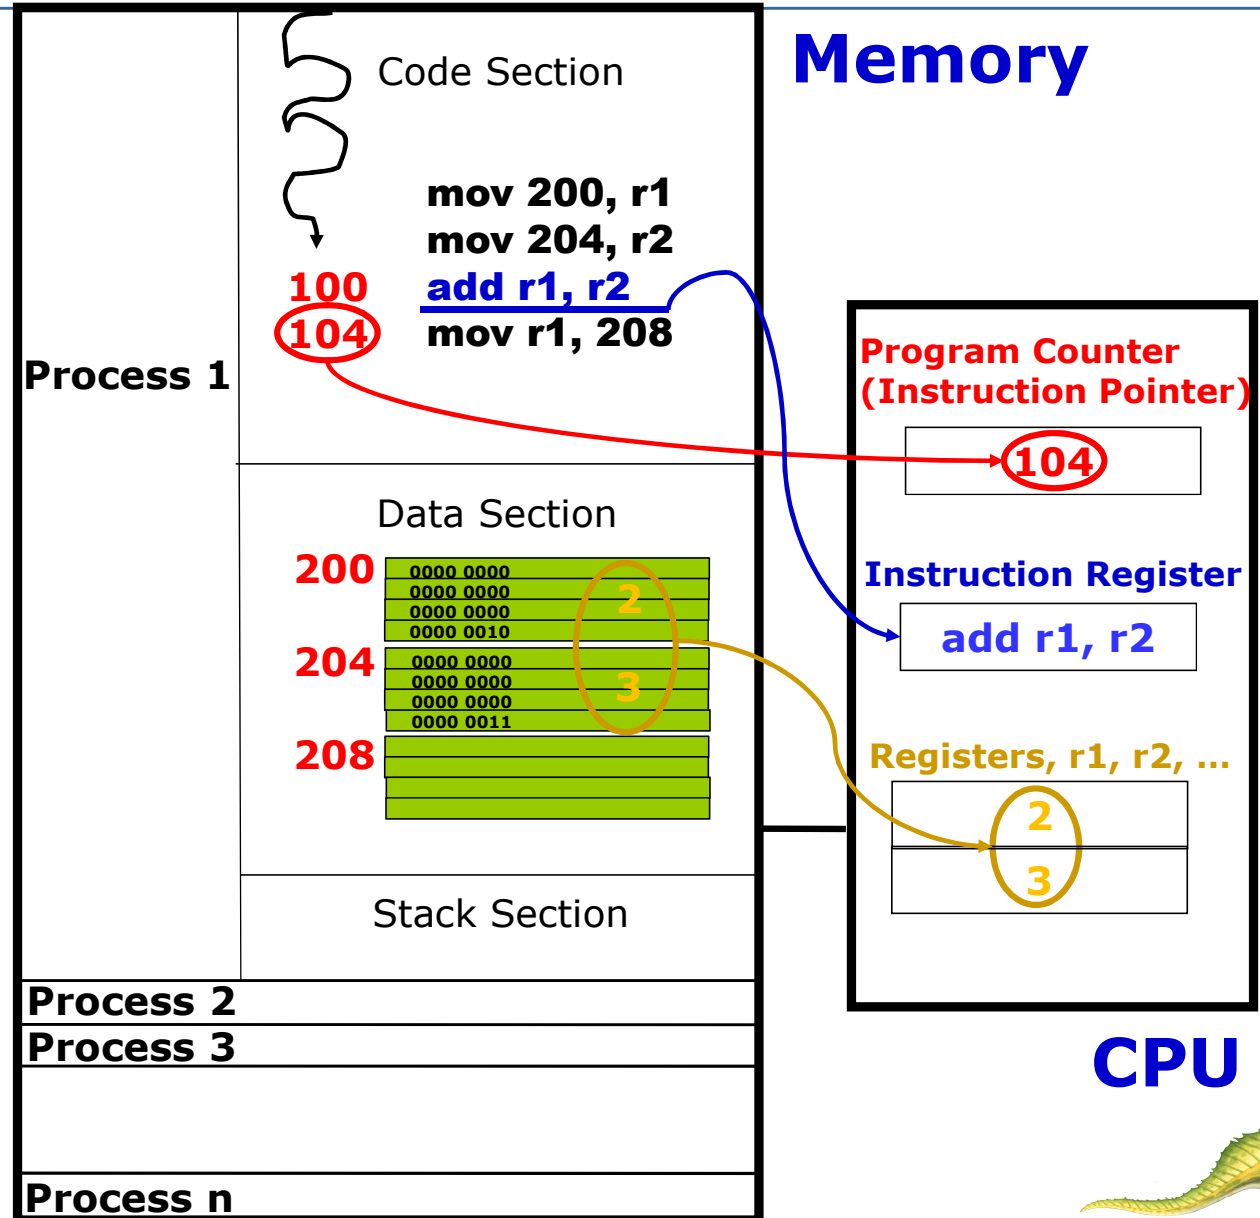
Assembly Language

```
i => address 200  
j => address 204  
k => address 208
```

```
mov 200, r1  
mov 204, r2  
add r1, r2  
mov r1, 208
```

Machine Language

```
010101....00011  
010110....01010  
010101....00011  
010110....01010
```





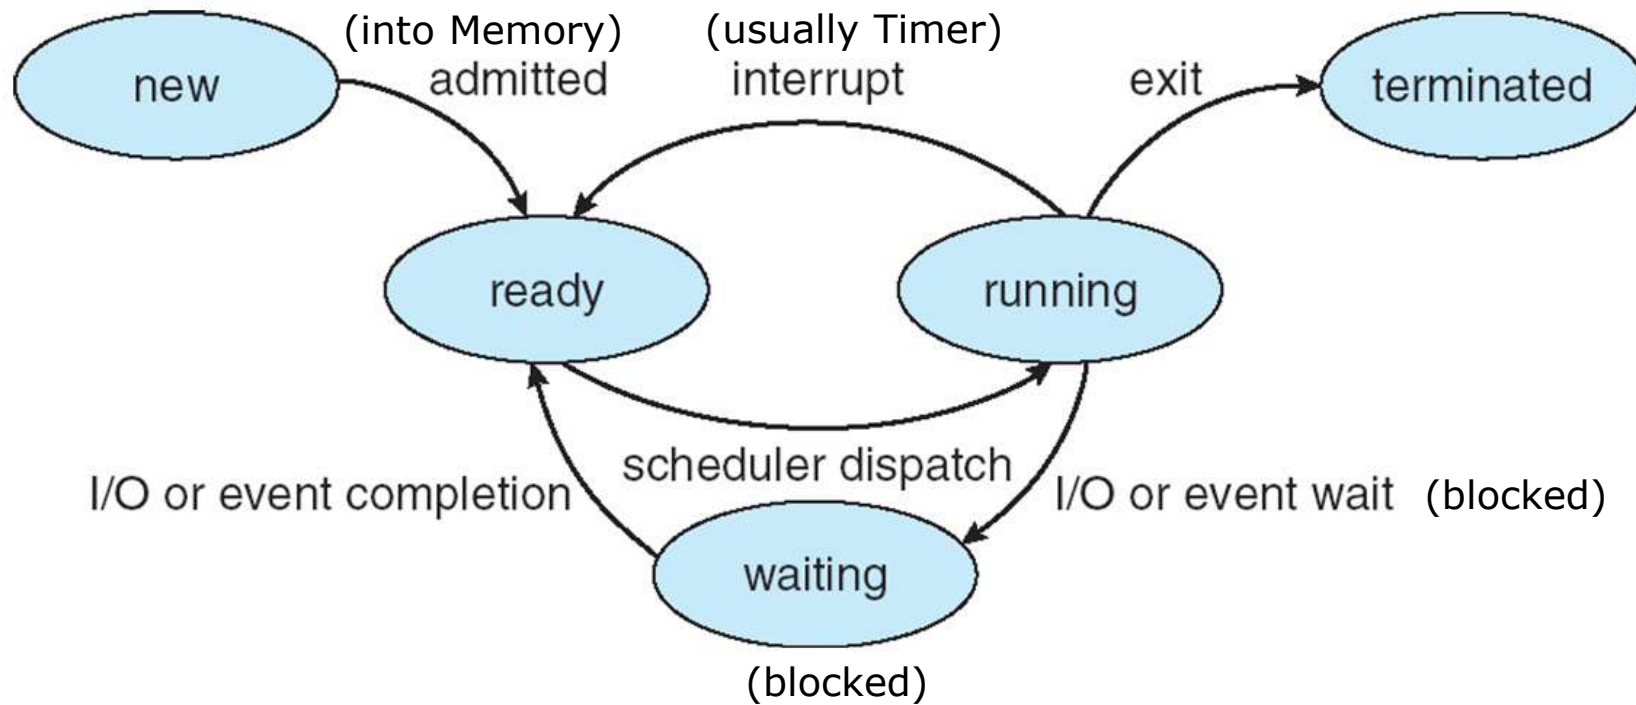
Process State

- As a process executes, it changes **state**
 - **new**: The process is being created
 - **ready**: The process is **waiting to be assigned to a processor**
 - **running**: **Instructions are being executed**
 - ▶ The only state that hold the CPU
 - **waiting(blocked)**: The process is **waiting(blocked) for some event to occur**
 - ▶ Cannot run on CPU
 - **terminated**: The process has finished execution





Diagram of Process State



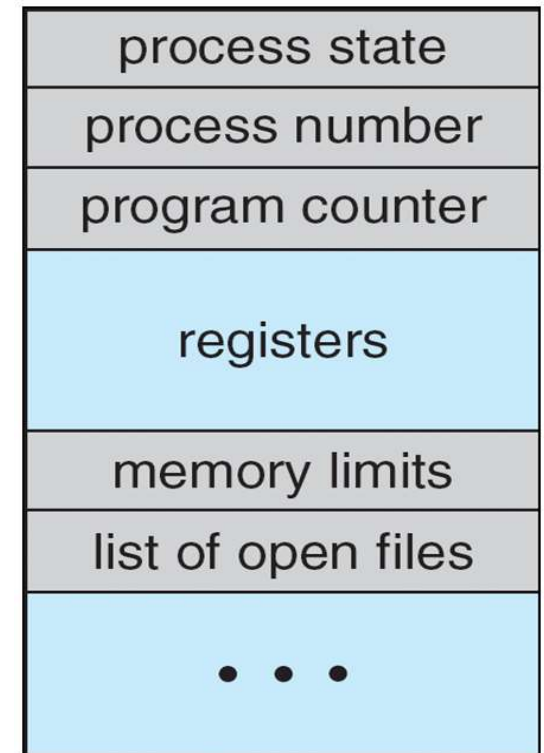


Process Control Block (PCB)

Information associated with each process

- Stored in main memory

- **Process state** – running, waiting, etc
- **Program counter** – location of instruction to next execute
- **CPU registers** – contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files





Process 1

PCB

Program Counter
104

Code Section

100

104

add r1, r2

200
300

2 Data Section
3

Stack Section

Process 2

PCB

Program Counter
2004

Code Section

2000

2004

Data Section

Stack Section

⋮

Process n

PCB

Program Counter
30004

Code Section

30000

30004

Data Section

Stack Section

Memory

CPU

Program Counter
(Instruction Pointer)

104

Instruction Register

add r1, r2

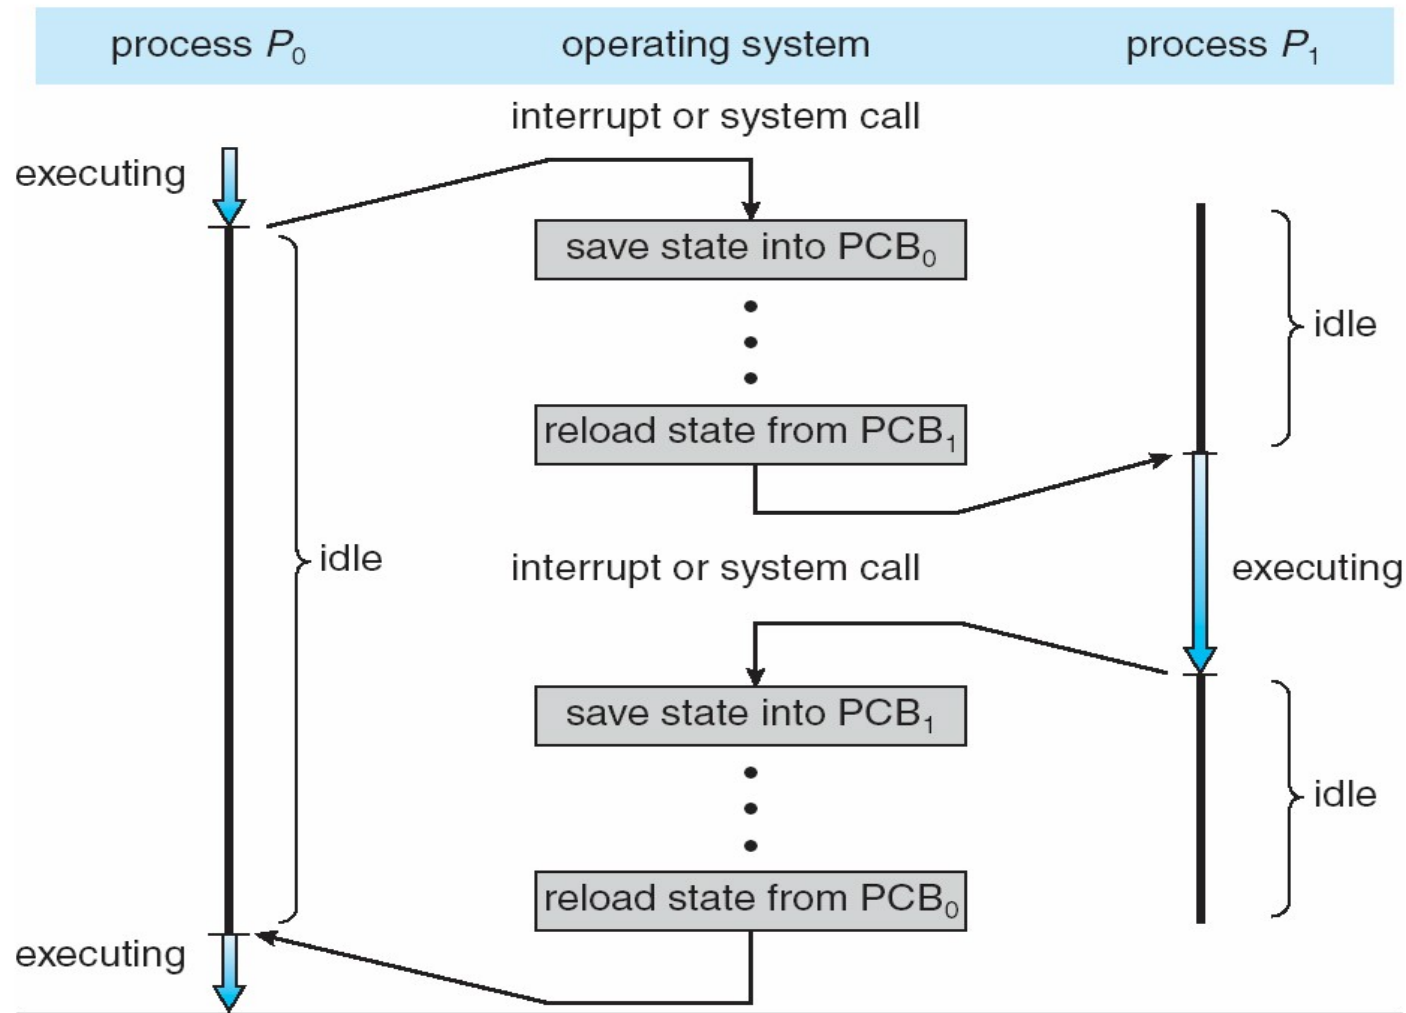
Registers, r1, r2, ...

2
3





CPU Switch From Process to Process





Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented **in the PCB**
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB -> longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU -> multiple contexts loaded at once





Operations on Processes

- System must provide mechanisms for process **creation**, **termination**, and so on as detailed next

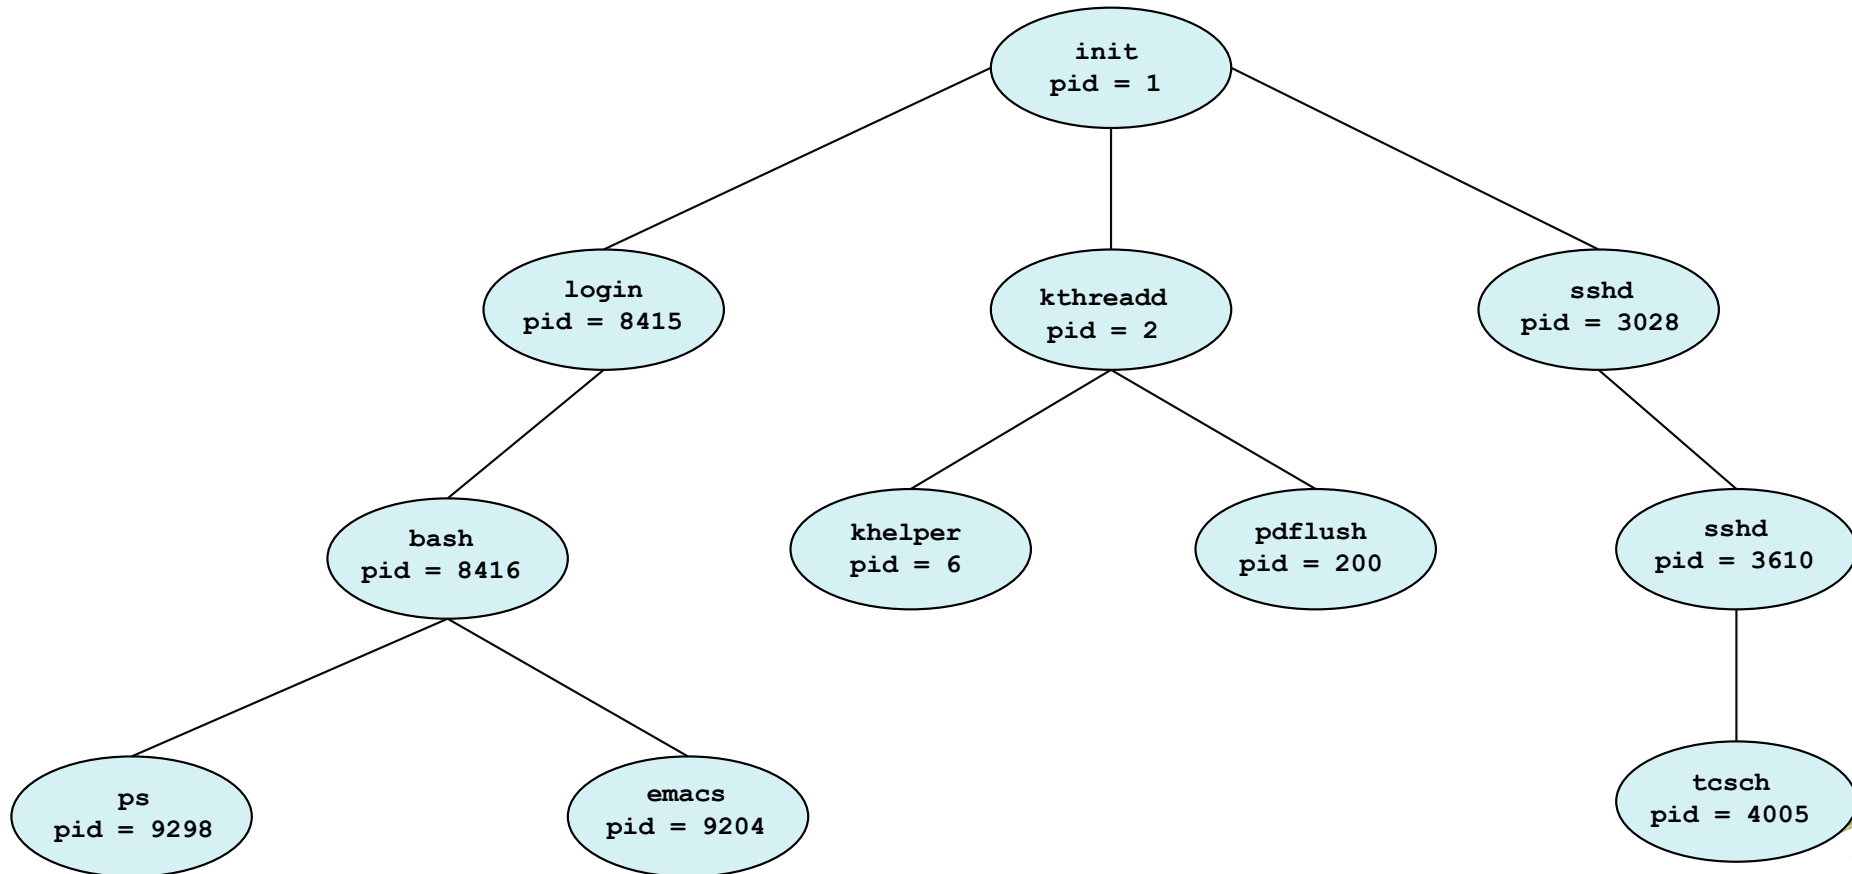
- Four Steps of **Process creation**
 - Create PCB within OS kernel
 - Allocate memory space
 - Load binary program
 - Initialization of program





Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- A Tree of Processes in Linux





Process Creation

- Resource(CPU time, memory, files, I/O devices) sharing options:
 - Parent and children share **all** resources
 - Children share **subset** of parent's resources
 - Parent and child share **no** resources
- Execution options:
 - Parent and children **execute concurrently**
 - **Parent waits** until some or all of its children have **terminated**
- **Address-space** possibilities for the new process:
 - Child process is **a duplicate of parent process** (it has the same program and data as the parent)
 - Child process has **a new program loaded into it**

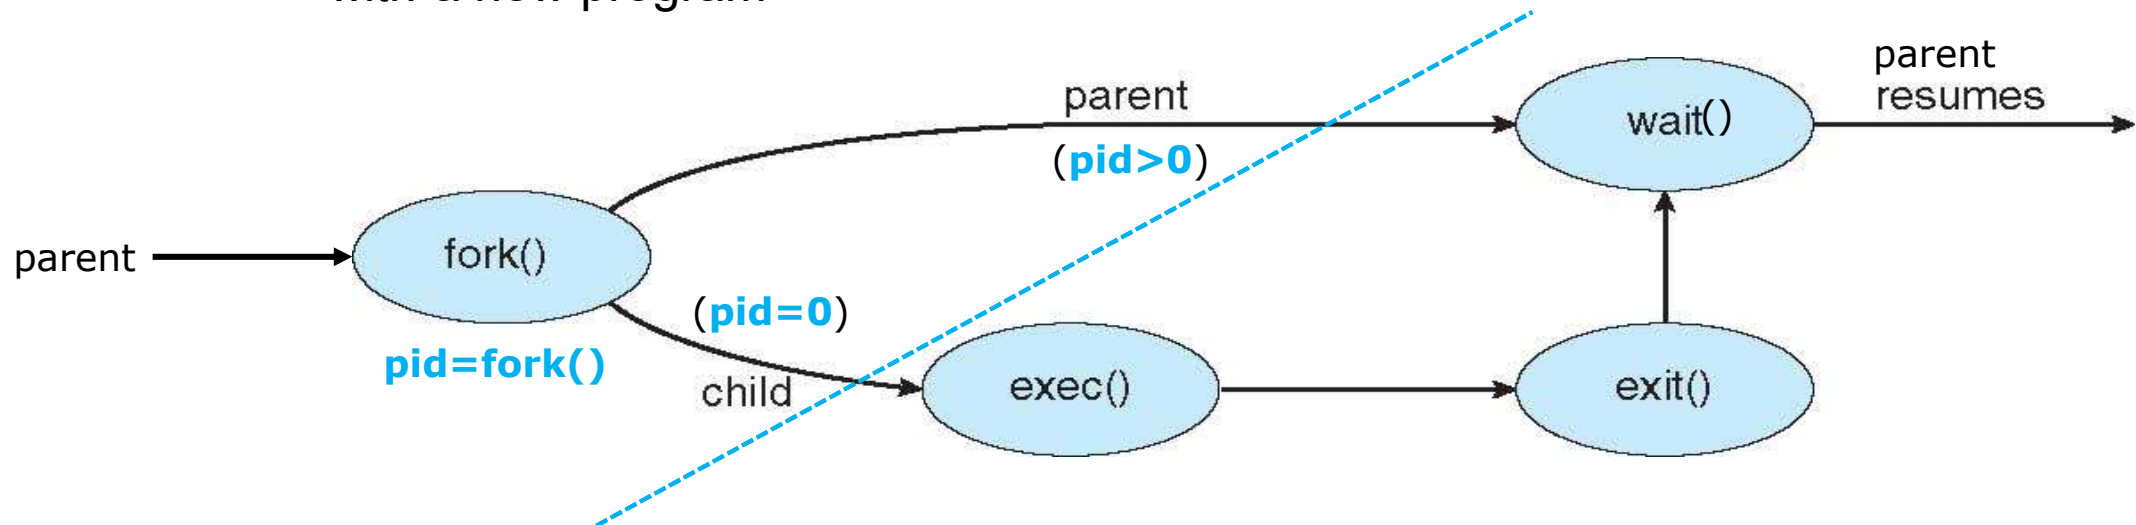




Process Creation (Cont.)

■ UNIX examples

- `fork()` system call creates new process
- `exec()` system call used after a `fork()` to *replace the process' memory space with* a new program





C Program Forking Separate Process in UNIX

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```





Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```





Process Termination

- Process terminates when it **executes last statement** and asks the OS to delete it by `exit()` system call
- Process may return a status value to its parent process via `wait()` system call
- Process' resources are deallocated by OS
- Parent process may wait for termination of a child process by using `wait()` :

```
pid t_pid; int status;  
pid = wait(&status);
```

- If no parent waiting, then terminated process is a **zombie**
- If parent terminated, processes are **orphans**





Multiprocess Architecture – Chrome Browser

- **Tabbed browsing** and **Active content**
 - A **single instance** of a **web browser** application may open **several websites at the same time**, each site in a separate tab
 - Website may contain **active content** such as JavaScript, Flash, HTML5 which may contain software bugs, resulting in **sluggish response times** and even **crash**
- If web browsers ran as a **single process**(some still do)
 - one web site causes trouble, entire browser can hang or crash





Multiprocess Architecture – Chrome Browser

- Google Chrome Browser is **multiprocess** with three different types
 - **Browser process** manages user interface, disk and network I/O
 - ▶ Created when Chrome is started; only one browser process is created
 - **Renderer process** renders web pages, deals with HTML, JavaScript, images
 - ▶ A new renderer process is created for each website opened in a new tab
 - **Plug-in process** for each type of plug-in such as Flash or QuickTime in use
- The **advantage** of multiprocess approach
 - **Websites run in isolation from one another**; if one website crashes, only its renderer process is affected





Interprocess Communication

- Processes within a system may be *independent* or *cooperating*
- *Independent process* cannot affect or be affected by the execution of another process
- *Cooperating process* can affect or be affected by the execution of another process
- **Reasons** for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **interprocess communication (IPC)**
 - Mechanism for processes to **communicate data** and to **synchronize their actions**





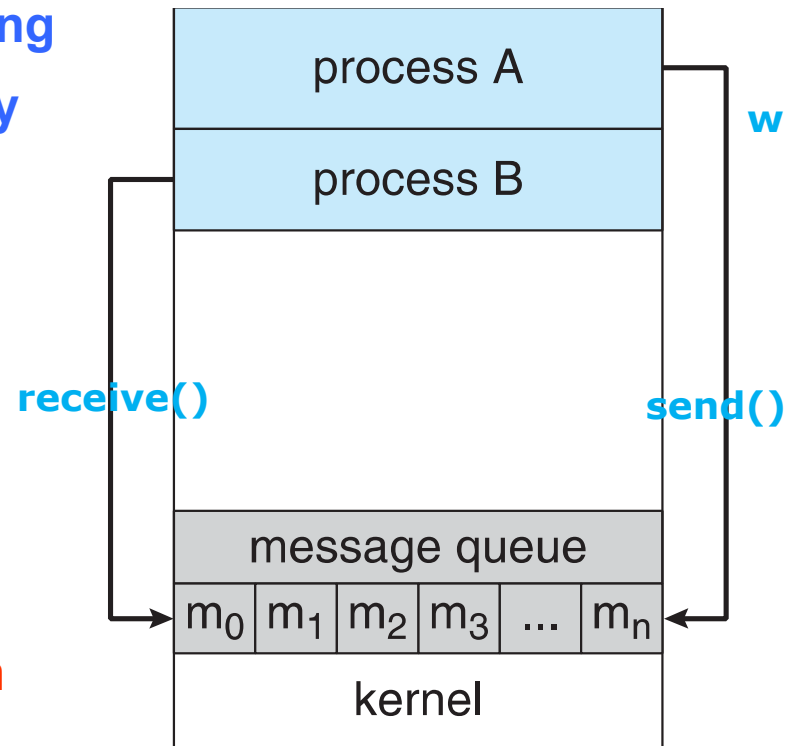
Interprocess Communication

■ Two Communications models of IPC

- Message passing
- Shared memory

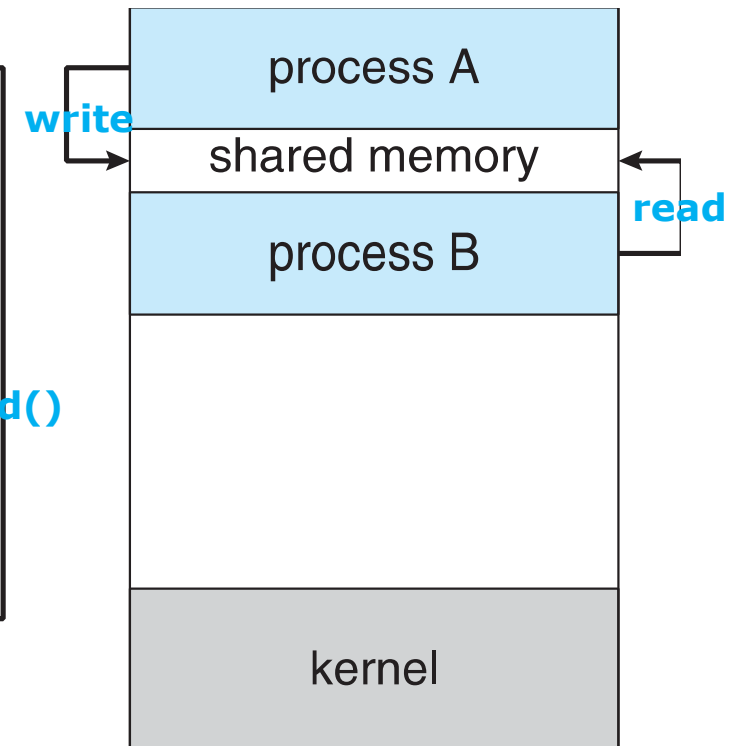
■ Synchronization

- Mutex
- Lock,
- Semaphore
- Critical Section
- See Ch.6



(a)

Message Passing



(b)

Shared Memory





Interprocess Communication – Message Passing

- Mechanism for processes to **communicate** and to **synchronize their actions**
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send(message)** – message size fixed or variable
 - **receive(message)**
- If P and Q wish to communicate, they need to:
 - establish a **communication link** between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., direct or indirect, synchronous or asynchronous, automatic or explicit **buffering**)





Direct Communication

- Processes must **name each other explicitly**:
 - **send** (*P*, *message*) – send a message to **process P**
 - **receive**(*Q*, *message*) – receive a message from **process Q**
- Properties of communication link
 - Links are established **automatically**
 - A link is associated with exactly **one pair of communicating processes**
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional





Indirect Communication

- Messages are directed and received from **mailboxes** (also referred to as **ports**)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox

- Properties of communication link
 - Link established only **if processes share a common mailbox**
 - A link may be associated **with many processes**
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional





Indirect Communication

■ Operations

- create a new mailbox
- send and receive messages through mailbox
- destroy a mailbox

■ Primitives are defined as:

send(A, message) – send a message to **mailbox A**

receive(A, message) – receive a message from **mailbox A**





Message Passing - Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue
 - **Non-blocking receive** has the receiver receive a valid message or null





Synchronization (Cont.)

- Different combinations possible
 - If **both send and receive are blocking**, we have a **rendezvous** between the sender and the receiver
- Producer-consumer problem becomes trivial

```
message next_produced;
while (true) {
    /* produce an item in next_produced */
    send(next_produced);
}
message next_consumed;
while (true) {
    receive(next_consumed);

    /* consume the item in next_consumed */
}
```





Buffering

- Queue of messages attached to the link; implemented in one of three ways
 1. **Zero capacity** – 0 messages
Sender must wait for receiver (rendezvous)
 2. **Bounded capacity** – finite length of n messages
Sender must wait if link full
 3. **Unbounded capacity** – infinite length
Sender never waits





Examples of IPC Systems - POSIX

■ POSIX Shared Memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

- Also used to open an existing segment to share it

- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Now the process could write to the shared memory

```
sprintf(shared memory, "Writing to shared memory");
```





```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
```

```
int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

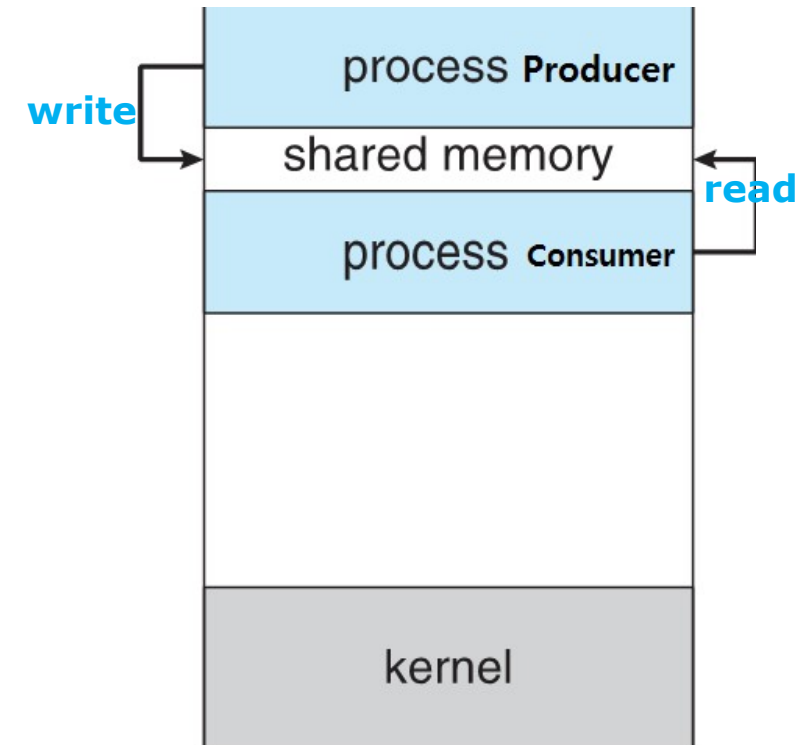
    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

IPC POSIX Producer





```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
```

```
int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

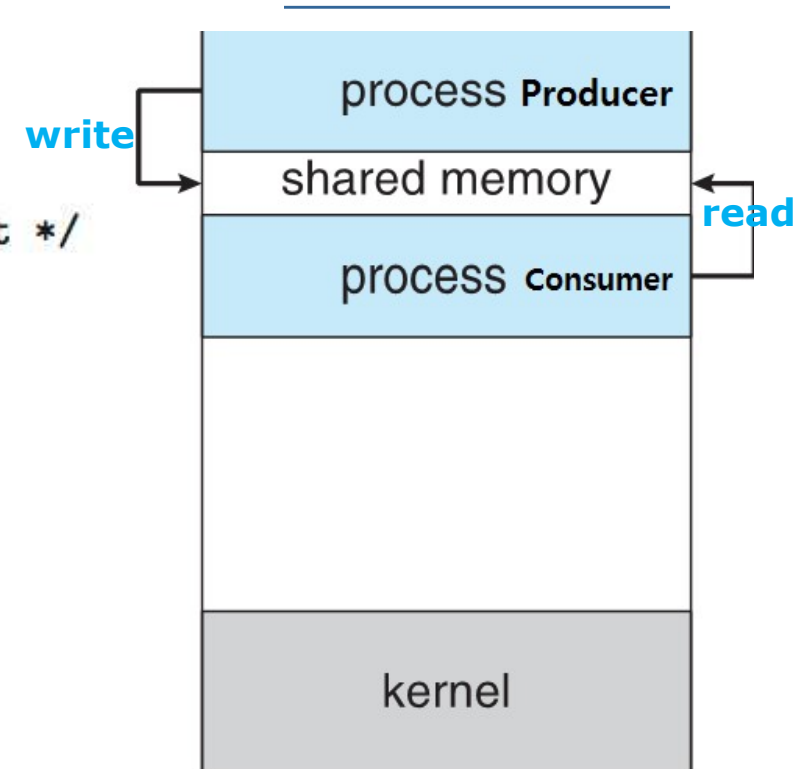
    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

IPC POSIX Consumer



End of Chapter 3

