


# Chapter 6: Synchronization

---





# Objectives

---

- To examine several **classical process-synchronization problems**
- To explore several tools that are used to solve process synchronization problems
  - **Windows API** for Process/Thread Synchronization : **Semaphores, Event, CRITICAL\_SECTION, Mutex**
  - **POSIX System Calls** for IPC(Interprocess Communication) : Pipes, FIFOs, Streams, **Semaphores, Message Queues, Shared Memory, Sockets**





# Background

---

- **Cooperating process** is one that can affect or be affected by other processes
- Cooperating processes can either
  - **directly share a logical address space**
    - ▶ **shared memory**, both **code** and **data** or
    - ▶ Through the use of **threads**
  - be allowed to **share data** through **files** or **messages**
- Processes can **execute concurrently**
  - May be **interrupted(context switching)** at any time, **partially completing execution**
- **Concurrent access to shared data** may result in **data inconsistency**
- **Maintaining data consistency** requires **mechanisms** to **ensure the orderly execution of cooperating processes**





# Background

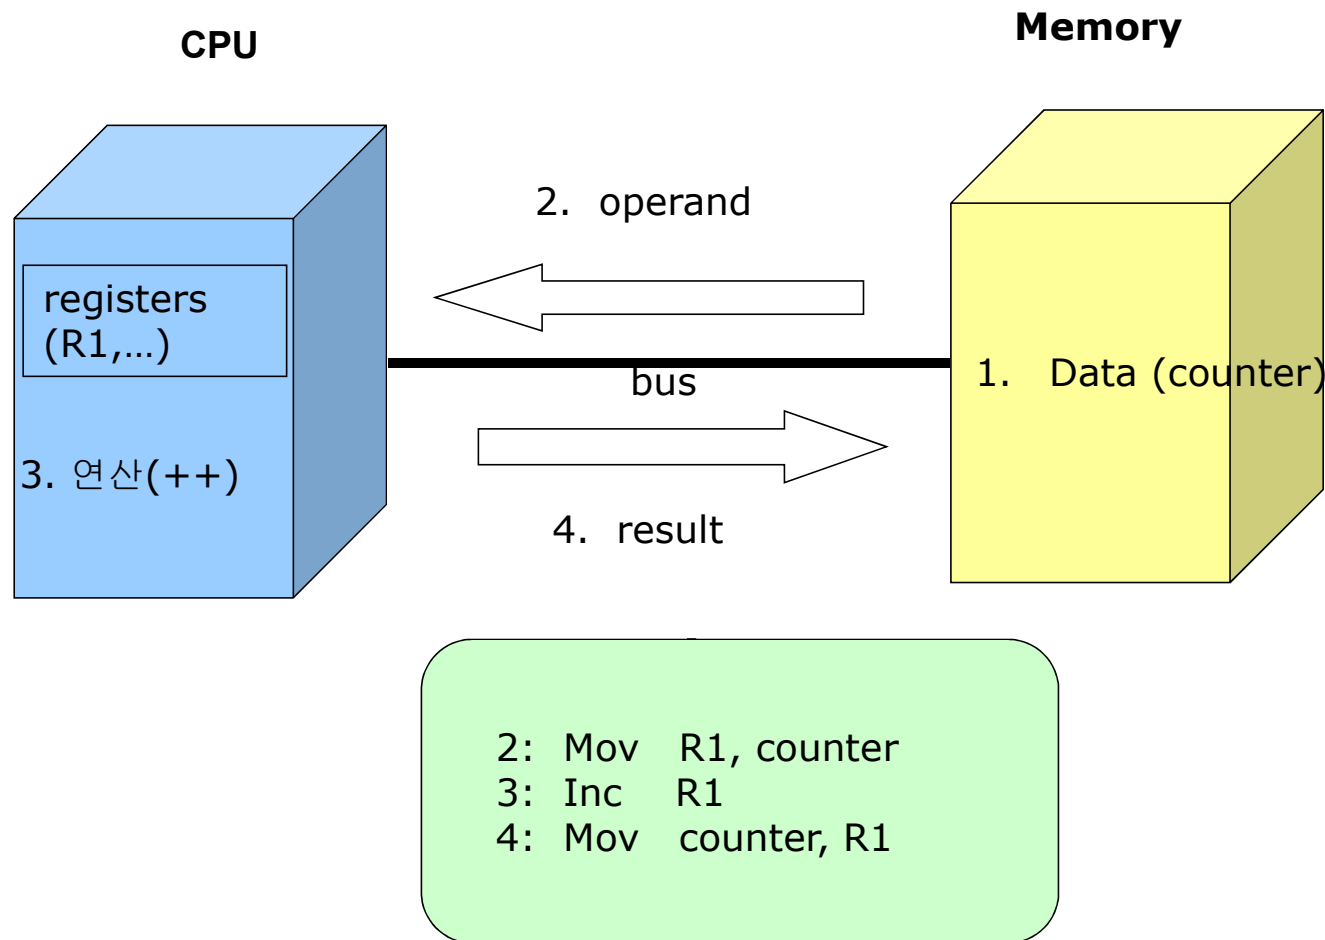
## ■ Race Condition

- The situation where several processes access and manipulate shared data concurrently.
  - The **final value** of the **shared data** depends upon **which process finishes last**.
- To prevent race conditions, **concurrent processes must be synchronized**.



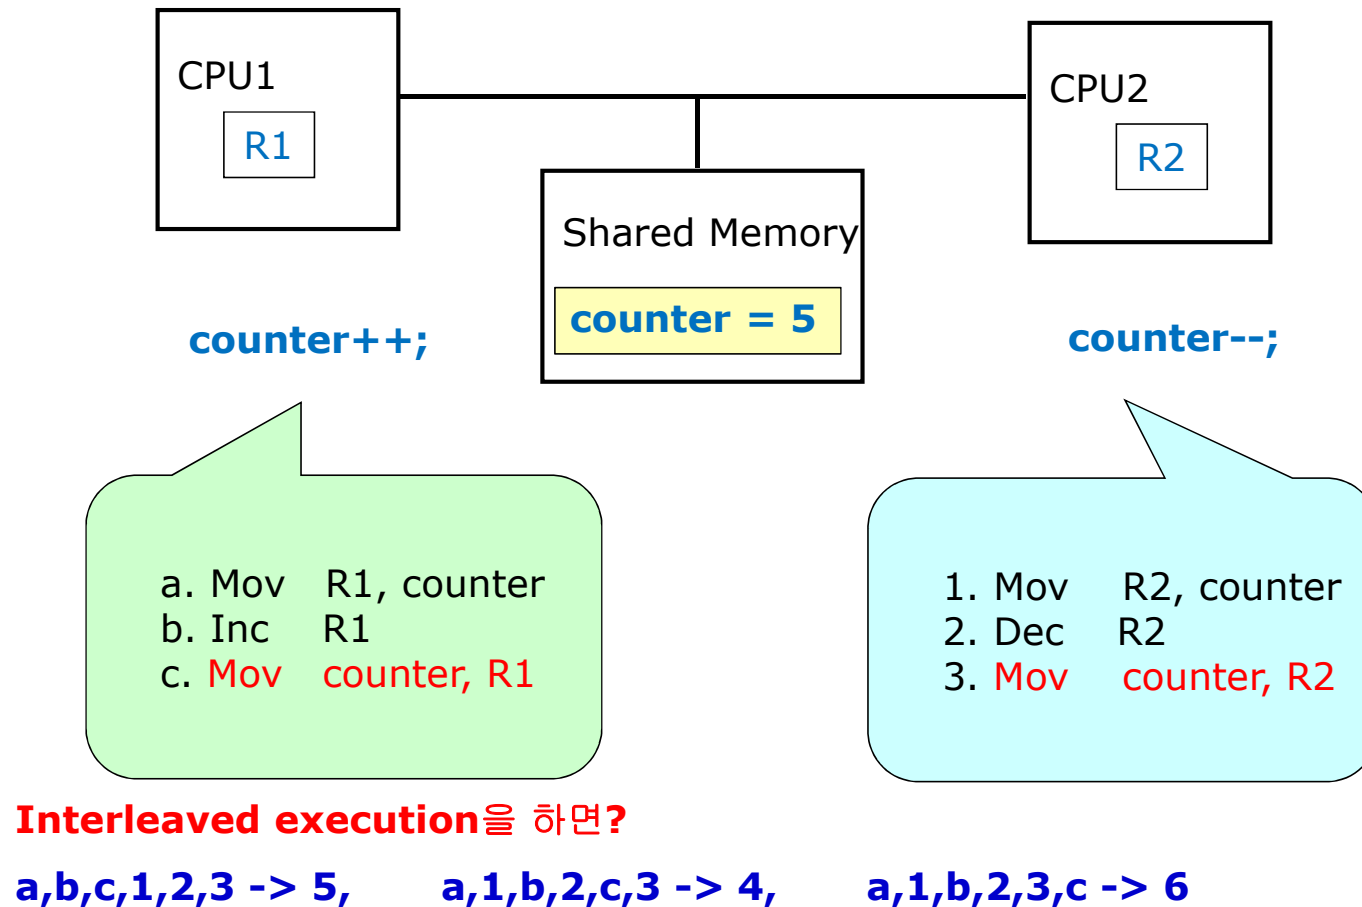


# “counter++” is not **ATOMIC**





# Example of a Race Condition



- We need to ensure that **only one process at a time can be manipulating the variable counter**
  - Serial component of Amdahl's Law





# Synchronization

## ■ Race Condition

A situation in which **multiple threads or processes** read and write the **same data concurrently** and the **outcome** of the execution depends on the particular order in which the access take place

## ■ Critical Section

A section of code within a process or a thread that requires **access to shared resources** and that may **not** be executed while another process is in a corresponding critical section.

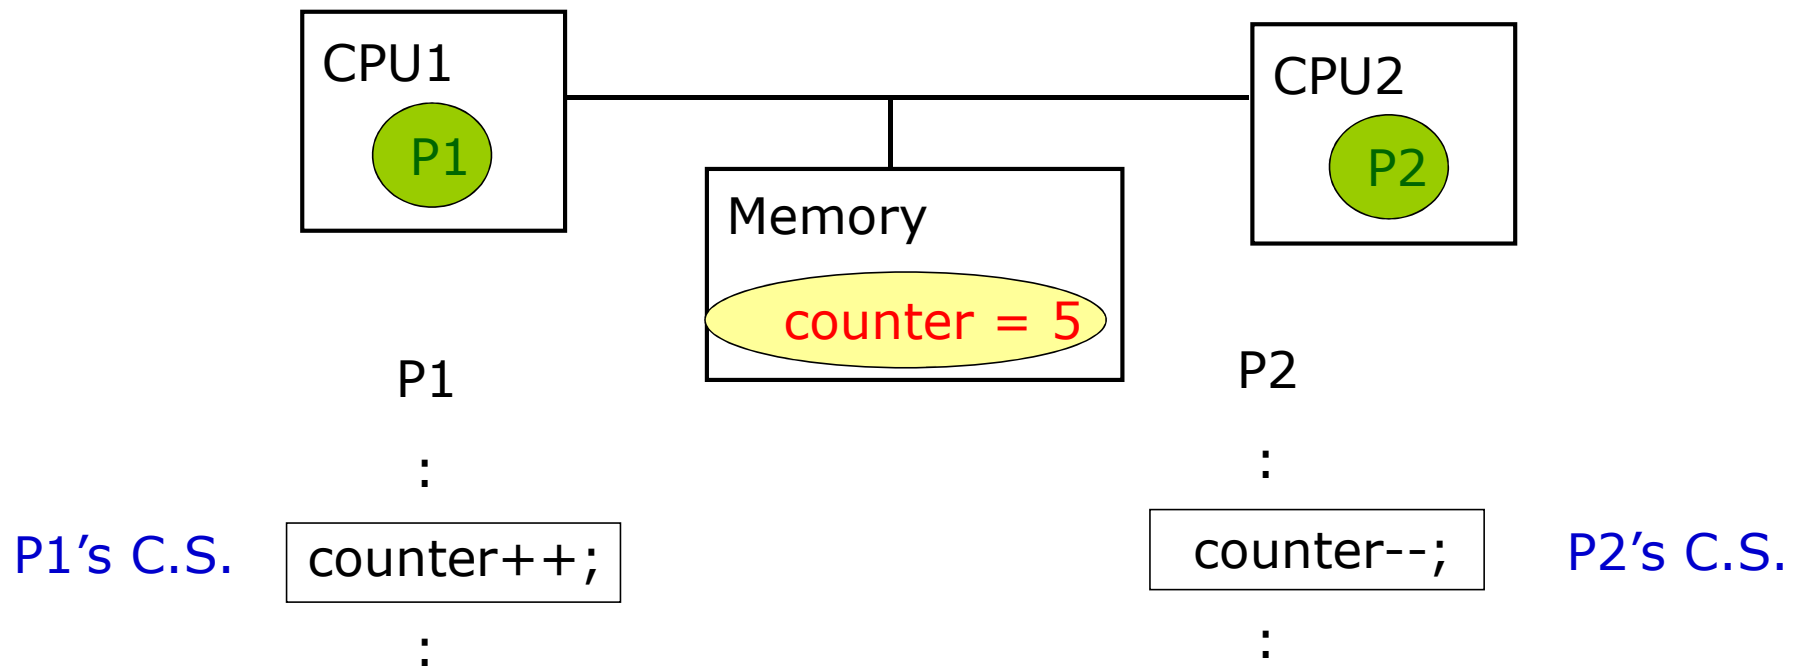
## ■ Mutual Exclusion

The requirement that when one process is in a critical section, no other process may be in a corresponding critical section.





# Critical Section



- **Maintaining data consistency** requires **mechanisms** to ensure the orderly execution of cooperating processes







# Mutex Locks

## ■ mutex lock

- **mutex** is short for **mut**ual **ex**clusion
- to protect critical regions and thus prevent race conditions
- A process must **acquire** the lock before entering a critical section
  - ▶ **acquire()** function
- It **releases** the lock when it exits the critical section
  - ▶ **release()** function
- a boolean variable **available** indicating if lock is available or not
- Calls to **acquire()** and **release()** must be **atomic**
  - ▶ Usually implemented via hardware atomic instructions

## ■ But this solution requires **busy waiting**

## ■ This lock therefore called a **spinlock**





# acquire() and release()

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
release() {  
    available = true;  
}
```





# Semaphores

- A more robust synchronization tool that provide more sophisticated ways for processes to synchronize their activities
- Semaphore S : integer variable
- Can only be accessed via **two indivisible (atomic) operations**
  - Originally called **P**(or **wait**) and **V**(or **signal**)

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal (S) {  
    S++;  
}
```

**P (S):** while ( $S \leq 0$ ) do *no-op*;  
          **S--;**

i.e. wait

If positive, decrement-&-enter.  
Otherwise, wait until it gets positive

**V (S):**  
          **S++;**





# Usage of Semaphore - Critical Section of $n$ Processes

- Shared data:

**semaphore  $s$** ; // shared semaphore among  $n$  processes  
// initialized to **1**

- Process  $P_i$ :

**repeat**

**wait( $s$ );** // entry section

**critical section**

**signal( $s$ );** // exit section

**remainder section**

**until false;**





# Usage of Semaphore - Synchronization of Two Processes

## ■ Problem

- Two concurrently running processes, P1 and P2
- The statement **S2 of P2** should be executed only **after** the statement **S1 of P1**.

## ■ Implementation

Shared data:

**Semaphore synch**; // P1 and P2 share a common semaphore  
// initialized to **0**

Process P1:

```
S1;  
signal(synch);
```

Process P2:

```
wait(synch);  
S2;
```





# Two Types of Semaphores

---

## ■ *Binary semaphore*

- integer value can range only between 0 and 1;

## ■ *Counting semaphore*

- integer value can range over an unrestricted domain.





# Classical Problems of Synchronization

- Classical problems used to test nearly every newly-proposed synchronization schemes
  - **Producer-consumer Problem(Bounded-Buffer Problem)**
    - ▶ 제한된 버퍼에 데이터를 채우고/가져가는 문제
  - **Readers and Writers Problem**
    - ▶ 데이터의 공유문제
  - **Dining-Philosophers Problem**
    - ▶ 한 자원을 가지고 다른 자원을 요청하는 문제





# Producer-Consumer Problem

---

- Paradigm for cooperating processes
- **producer** process produces **information** that is consumed by a **consumer** process
  - **unbounded-buffer** places no practical limit on the size of the buffer
  - **bounded-buffer** assumes that there is a fixed buffer size








# Bounded-Buffer – Shared-Memory Solution

## ■ Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```


## ■ Solution is correct, but can only use BUFFER\_SIZE-1 elements





# Producer - Busy Waiting Solution


```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```





# Consumer - Busy Waiting Solution


```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
  
    /* consume the item in next_consumed */  
}
```





# Producer-Consumer Problem : Busy Waiting Solution

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER SIZE) ;  
        /* do nothing */  
    buffer[in] = next produced;  
    in = (in + 1) % BUFFER SIZE; counter++;  
}
```




**buffer FULL**

## Producer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next consumed = buffer[out];  
    out = (out + 1) % BUFFER SIZE; counter--;  
    /* consume the item in next consumed */  
}
```

## Consumer<sub>in</sub>



**buffer EMPTY**





# Producer-Consumer Problem : Semaphore Solution

---

- $n$  buffers, each can hold one item
- Semaphore `mutex` initialized to the value 1
- Semaphore `full` initialized to the value 0
- Semaphore `empty` initialized to the value  $n$





# Producer – Semaphore Solution

- The structure of the **producer** process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```





# Consumer – Semaphore Solution

- The structure of the **consumer** process

```
do {  
    wait(full) ;  
    wait(mutex) ;  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex) ;  
    signal(empty) ;  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true) ;
```






# Producer-Consumer Problem : Semaphore Solution

## Shared data

semaphore **full** = 0, **empty** = n, **mutex** = 1;



## Producer

```
do {  
    ...  
    produce an item  
    ...  
    wait(empty); // P(empty)  
    wait(mutex);  
    ...  
    add item to buffer  
    ...  
    signal(mutex);  
    signal(full); // V(full)  
} while (1);
```

## Consumer

```
do {  
    wait(full) // P(full)  
    wait(mutex);  
    ...  
    remove an item from buffer  
    ...  
    signal(mutex);  
    signal(empty); // V(empty)  
    ...  
    consume the item  
    ...  
} while (1);
```

