


Chapter 6: Synchronization





Objectives

- To examine several **classical process-synchronization problems**
- To explore several tools that are used to solve process synchronization problems
 - **Windows API** for Process/Thread Synchronization : **Semaphores, Event, CRITICAL_SECTION, Mutex**
 - **POSIX System Calls** for IPC(Interprocess Communication) : Pipes, FIFOs, Streams, **Semaphores, Message Queues, Shared Memory, Sockets**





Background

- **Cooperating process** is one that can affect or be affected by other processes
- Cooperating processes can either
 - **directly share a logical address space**
 - ▶ **shared memory**, both **code** and **data** or
 - ▶ Through the use of **threads**
 - be allowed to **share data** through **files** or **messages**
- Processes can **execute concurrently**
 - May be **interrupted(context switching)** at any time, **partially completing execution**
- **Concurrent access to shared data** may result in **data inconsistency**
- **Maintaining data consistency** requires **mechanisms** to **ensure the orderly execution of cooperating processes**





Background

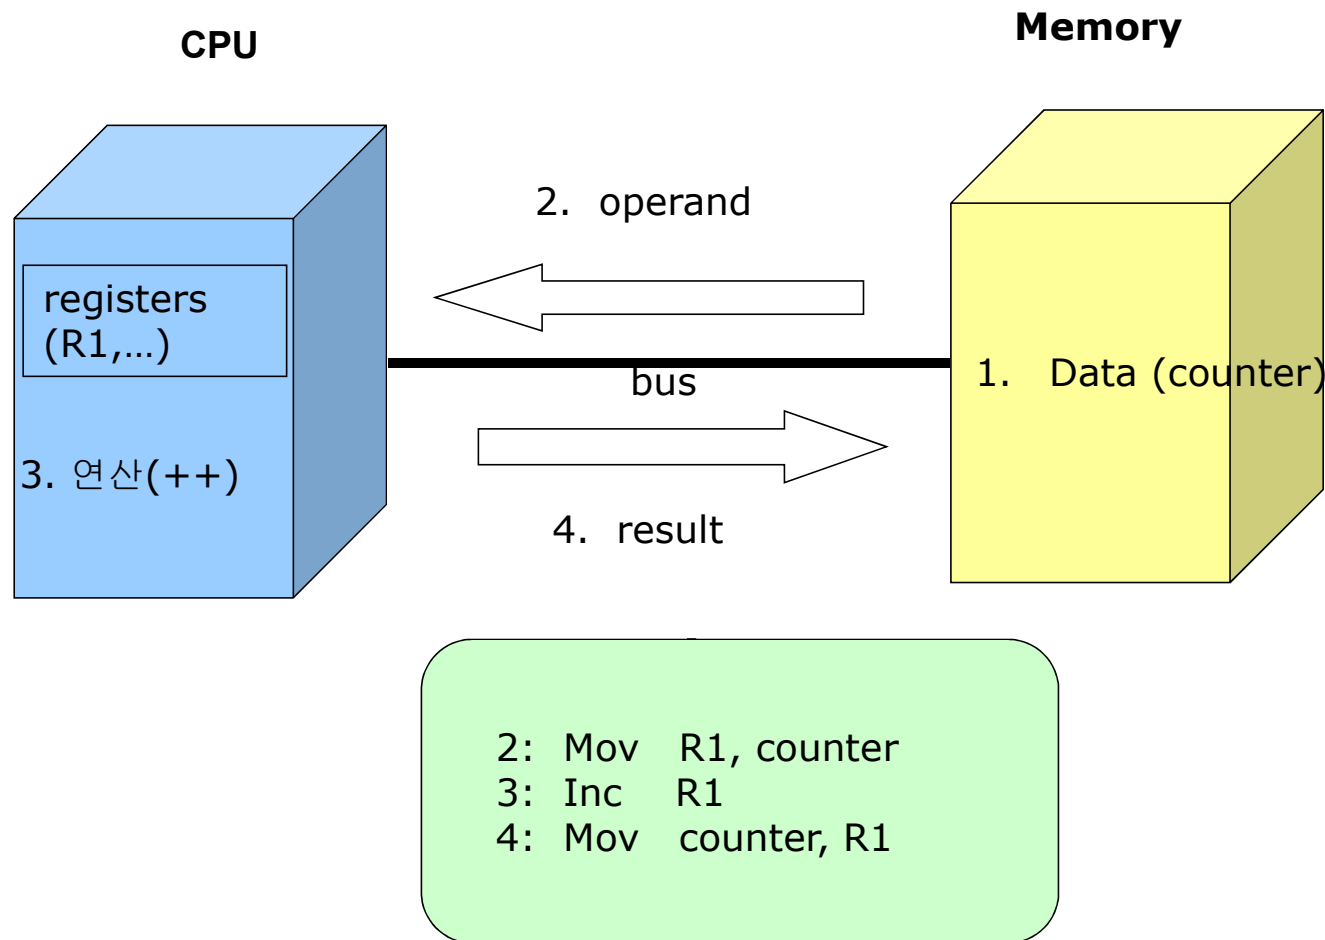
■ Race Condition

- The situation where several processes access and manipulate shared data concurrently.
 - The **final value** of the **shared data** depends upon **which process finishes last**.
- To prevent race conditions, **concurrent processes must be synchronized**.



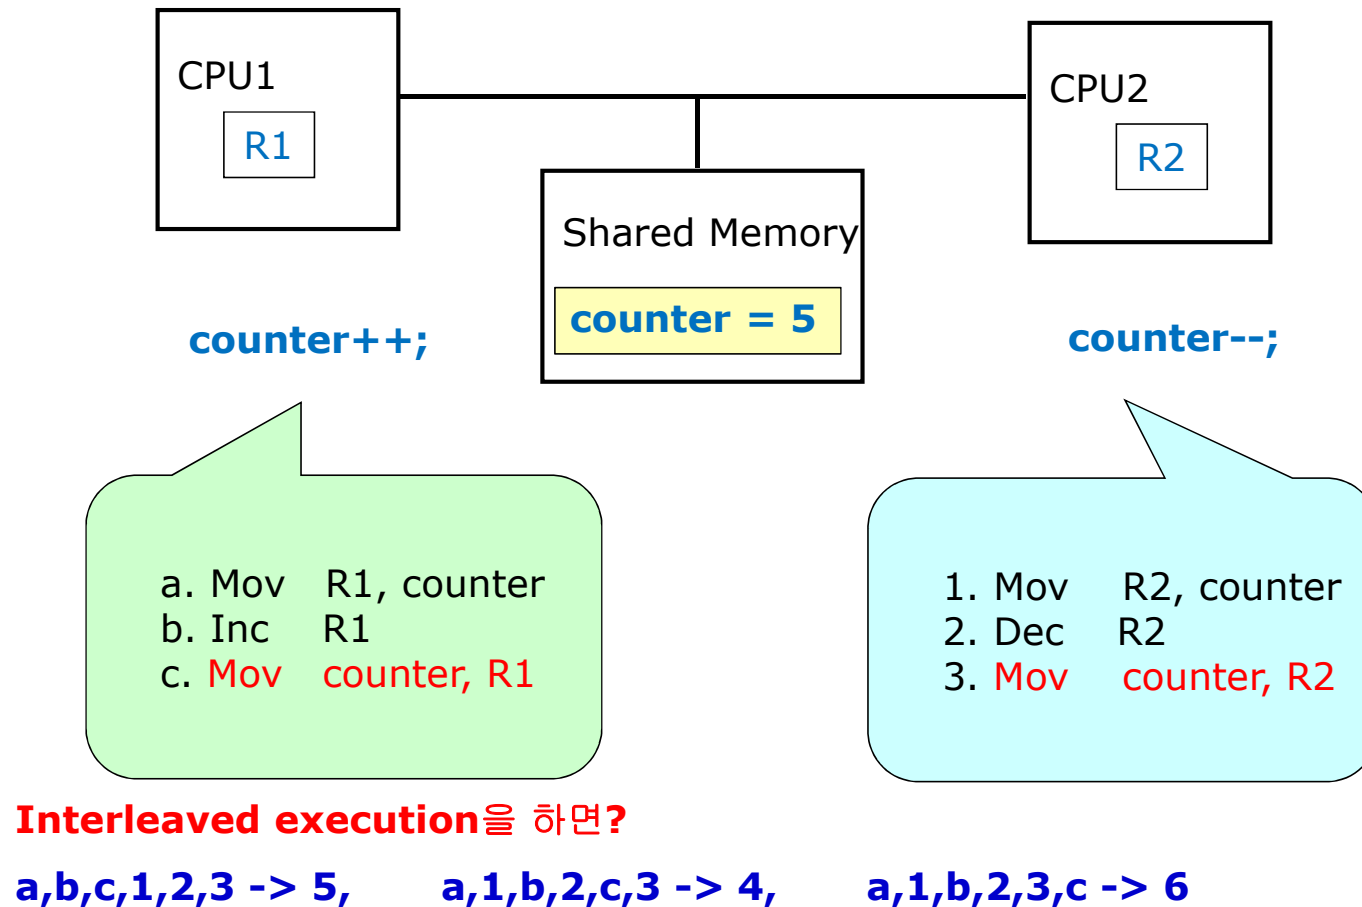


“counter++” is not **ATOMIC**





Example of a Race Condition



- We need to ensure that **only one process at a time can be manipulating the variable counter**
 - Serial component of Amdahl's Law





Synchronization

■ Race Condition

A situation in which **multiple threads or processes** read and write the **same data concurrently** and the **outcome** of the execution depends on the particular order in which the access take place

■ Critical Section

A section of code within a process or a thread that requires **access to shared resources** and that may **not** be executed while another process is in a corresponding critical section.

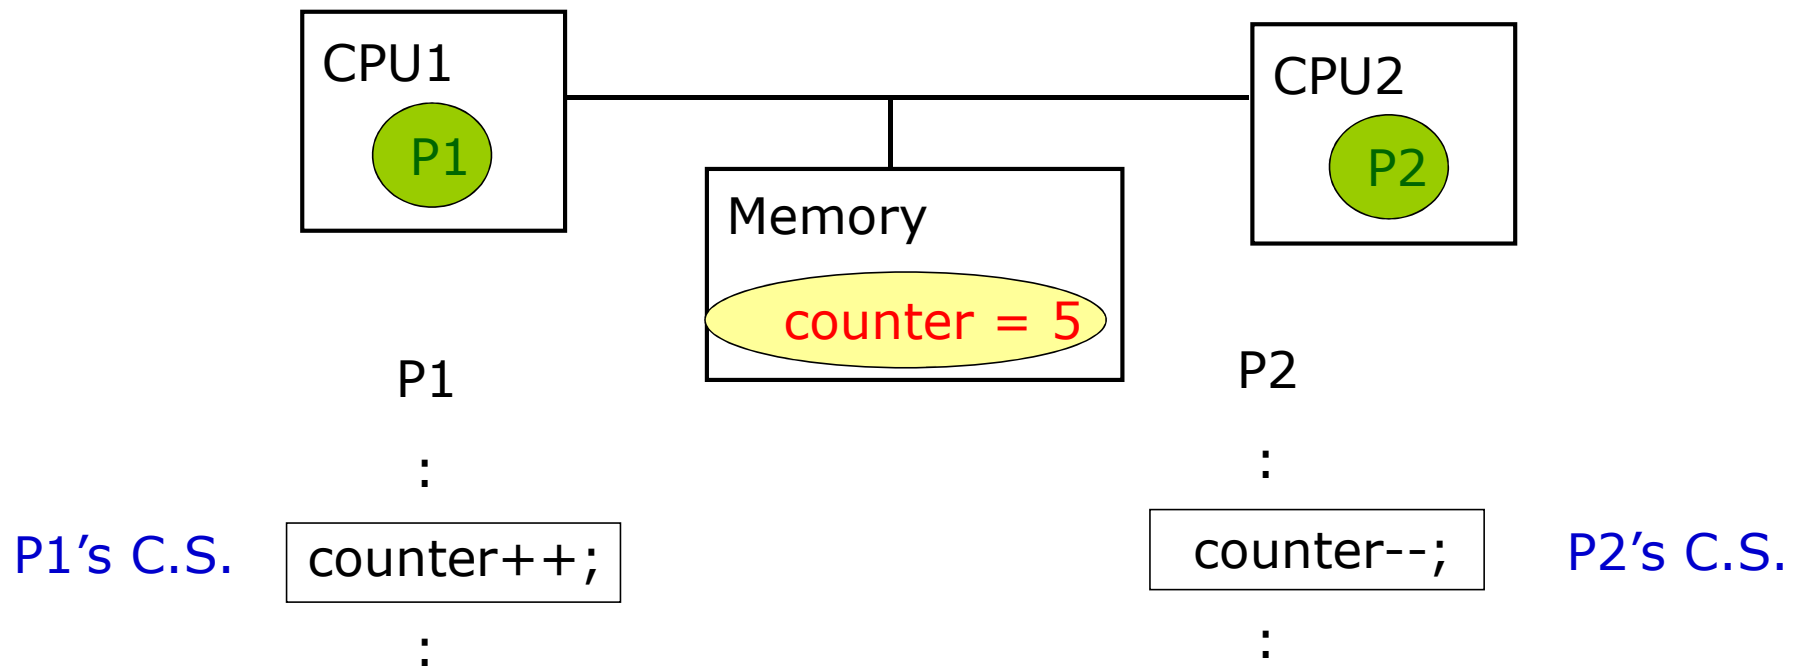
■ Mutual Exclusion

The requirement that when one process is in a critical section, no other process may be in a corresponding critical section.





Critical Section



- **Maintaining data consistency** requires mechanisms to ensure the orderly execution of cooperating processes





Mutex Locks

■ mutex lock

- **mutex** is short for **mut**ual **ex**clusion
- to protect critical regions and thus prevent race conditions
- A process must **acquire** the lock before entering a critical section
 - ▶ **acquire()** function
- It **releases** the lock when it exits the critical section
 - ▶ **release()** function
- a boolean variable **available** indicating if lock is available or not
- Calls to **acquire()** and **release()** must be **atomic**
 - ▶ Usually implemented via hardware atomic instructions

■ But this solution requires **busy waiting**

■ This lock therefore called a **spinlock**





acquire() and release()

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
release() {  
    available = true;  
}
```





Semaphores

- A more robust synchronization tool that provide more sophisticated ways for processes to synchronize their activities
- Semaphore S : integer variable
- Can only be accessed via **two indivisible (atomic) operations**
 - Originally called **P**(or **wait**) and **V**(or **signal**)

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

```
signal (S) {  
    S++;  
}
```

P (S): while ($S \leq 0$) do *no-op*;
 S--;

i.e. wait

If positive, decrement-&-enter.
Otherwise, wait until it gets positive

V (S):
 S++;





Usage of Semaphore - Critical Section of n Processes

- Shared data:

semaphore s ; // shared semaphore among n processes
// initialized to **1**

- Process P_i :

repeat

wait(s); // entry section

critical section

signal(s); // exit section

remainder section

until false;





Usage of Semaphore - Synchronization of Two Processes

- Problem
 - Two concurrently running processes, P1 and P2
 - The statement **S2 of P2** should be executed only **after** the statement **S1 of P1**.
- Implementation

Shared data:

Semaphore synch; // P1 and P2 share a common semaphore
// initialized to **0**

Process P1:

```
S1;  
signal(synch);
```

Process P2:

```
wait(synch);  
S2;
```





Two Types of Semaphores

■ *Binary semaphore*

- integer value can range only between 0 and 1;

■ *Counting semaphore*

- integer value can range over an unrestricted domain.





Classical Problems of Synchronization

- Classical problems used to test nearly every newly-proposed synchronization schemes
 - **Producer-consumer Problem(Bounded-Buffer Problem)**
 - ▶ 제한된 버퍼에 데이터를 채우고/가져가는 문제
 - **Readers and Writers Problem**
 - ▶ 데이터의 공유문제
 - **Dining-Philosophers Problem**
 - ▶ 한 자원을 가지고 다른 자원을 요청하는 문제





Producer-Consumer Problem

- Paradigm for cooperating processes
- **producer** process produces **information** that is consumed by a **consumer** process
 - **unbounded-buffer** places no practical limit on the size of the buffer
 - **bounded-buffer** assumes that there is a fixed buffer size






Bounded-Buffer – Shared-Memory Solution

■ Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```


■ Solution is correct, but can only use BUFFER_SIZE-1 elements





Producer - Busy Waiting Solution


```
item next_produced;
while (true) {
    /* produce an item in next_produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```





Consumer - Busy Waiting Solution


```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
  
    /* consume the item in next_consumed */  
}
```





Producer-Consumer Problem : Busy Waiting Solution

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER SIZE) ;  
        /* do nothing */  
    buffer[in] = next produced;  
    in = (in + 1) % BUFFER SIZE; counter++;  
}
```




buffer FULL

Producer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next consumed = buffer[out];  
    out = (out + 1) % BUFFER SIZE; counter--;  
    /* consume the item in next consumed */  
}
```

Consumer_{in}



buffer EMPTY





Producer-Consumer Problem : Semaphore Solution

- n buffers, each can hold one item
- Semaphore `mutex` initialized to the value 1
- Semaphore `full` initialized to the value 0
- Semaphore `empty` initialized to the value n





Producer – Semaphore Solution

- The structure of the **producer** process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```





Consumer – Semaphore Solution

- The structure of the **consumer** process

```
do {  
    wait(full) ;  
    wait(mutex) ;  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex) ;  
    signal(empty) ;  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true) ;
```






Producer-Consumer Problem : Semaphore Solution

Shared data

semaphore **full** = 0, **empty** = n, **mutex** = 1;



Producer

```
do {  
    ...  
    produce an item  
    ...  
    wait(empty); // P(empty)  
    wait(mutex);  
    ...  
    add item to buffer  
    ...  
    signal(mutex);  
    signal(full); // V(full)  
} while (1);
```

Consumer

```
do {  
    wait(full) // P(full)  
    wait(mutex);  
    ...  
    remove an item from buffer  
    ...  
    signal(mutex);  
    signal(empty); // V(empty)  
    ...  
    consume the item  
    ...  
} while (1);
```

